

MOVIE RECOMMENDER SYSTEM

Introduction

This project focuses on building a movie recommender system based on the Movielens dataset. It consists of 100,000 ratings by 610 unique users for 9724 unique movies. The objective is to build a system that can recommend top 5 movies to the users based on their ratings using collaborative filtering. This project addresses the business problem of personalized movie recommendations using the Movielens dataset. The dataset contains user ratings, movie metadata (title, genres, IDs), and tags, making it well-suited for collaborative filtering approaches since it captures both user behavior and movie attributes. The final dataset retained userId, movieId, title, genres, rating, and rating year, which represent the most important features for capturing user preferences over time.

For data preparation, we dropped duplicate identifiers, merged multiple files (ratings, movies, links), and imputed missing values with zeroes to ensure all users and movies were represented in the interaction matrix. We also applied scaling for model stability. Data preparation was handled with pandas (merging/cleaning), scikit-learn (SimpleImputer, StandardScaler, Pipeline), and numpy (matrix operations).

Modeling was performed using collaborative filtering via matrix factorization. Specifically, we built a Pipeline with TruncatedSVD for dimensionality reduction and Ridge Regression as the estimator. Hyperparameters were tuned with GridSearchCV, enabling cross-validated selection of latent factor regularization strength. Evaluation was conducted with train-test split and 5-fold cross-validation. The model achieved strong predictive performance, with RMSE of 3.29.

Import libraries

```
In [ ]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.pipeline import Pipeline
from sklearn.decomposition import TruncatedSVD
from sklearn.linear_model import Ridge
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error, make_scorer
import warnings
warnings.filterwarnings('ignore')
```

Load and merge datasets

```
In [2]: # Load and display the first few rows of the dataset
df_links=pd.read_csv('ml-latest-small\links.csv')
df_links.head()
```

```
Out[2]:
```

	movieId	imdbId	tmdbId
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

```
In [3]: #Load and display the first few rows of the movies dataset
df_movies=pd.read_csv('ml-latest-small\movies.csv')
df_movies.head()
```

```
Out[3]:
```

	movieId	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

```
In [4]: #Load and display the first few rows of the ratings dataset
df_ratings=pd.read_csv(r'ml-latest-small\ratings.csv')
df_ratings.head()
```

```
Out[4]:
```

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931

```
In [5]: #Load and display the first few rows of the tags dataset
df_tags=pd.read_csv(r'ml-latest-small\tags.csv')
df_tags.head()
```

```
Out[5]:
```

	userId	movieId	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

```
In [8]: #Merge all datasets into a single dataframe
ratings_tags = pd.merge(df_ratings, df_tags, on=["userId", "movieId"], how="left")
ratings_tags_movies = pd.merge(ratings_tags, df_movies, on="movieId", how="left")
df_merged = pd.merge(ratings_tags_movies, df_links, on="movieId", how="left")
df_merged.head()
```

Out[8]:	userId	movieId	rating	timestamp_x	tag	timestamp_y	title	genres	imdbId
0	1	1	4.0	964982703	NaN	NaN	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	114709
1	1	3	4.0	964981247	NaN	NaN	Grumpier Old Men (1995)	Comedy Romance	113228
2	1	6	4.0	964982224	NaN	NaN	Heat (1995)	Action Crime Thriller	113277
3	1	47	5.0	964983815	NaN	NaN	Seven (a.k.a. Se7en) (1995)	Mystery Thriller	114369
4	1	50	5.0	964982931	NaN	NaN	Usual Suspects, The (1995)	Crime Mystery Thriller	114814

```
In [7]: # Display information about the final dataframe
df_merged.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 102677 entries, 0 to 102676
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   userId          102677 non-null  int64
1   movieId         102677 non-null  int64
2   rating          102677 non-null  float64
3   timestamp_x     102677 non-null  int64
4   tag             3476 non-null   object
5   timestamp_y     3476 non-null   float64
6   title           102677 non-null  object
7   genres          102677 non-null  object
8   imdbId          102677 non-null  int64
9   tmdbId          102664 non-null  float64
dtypes: float64(3), int64(4), object(3)
memory usage: 7.8+ MB
```

Data Cleaning

```
In [9]: # Cleaninng and renaming columns for clarity
df_merged = df_merged.rename(columns={
    "timestamp_x": "rating_timestamp",
    "timestamp_y": "tag_timestamp"
})
```

```
In [10]: # Checking for duplicates
df_merged.duplicated().sum()
```

Out[10]: 0

Exploratory Data Analysis

```
In [11]: # Checking for missing values
df_merged.isnull().sum()
```


```
Out[11]:  userId          0
         movieId        0
         rating         0
         rating_timestamp 0
         tag           99201
         tag_timestamp   99201
         title          0
         genres         0
         imdbId         0
         tmdbId         13
         dtype: int64
```

```
In [12]: # Unique users and movies
n_users = df_merged['userId'].nunique()
n_movies = df_merged['movieId'].nunique()
print(f"\nUnique users: {n_users}, Unique movies: {n_movies}")
```

Unique users: 610, Unique movies: 9724

```
In [13]: # Checking rating usefulness; if ratings are varied and not all the same, they c
df_merged['rating'].describe()

plt.figure(figsize=(6,4))
df_merged['rating'].hist(bins=10)
plt.title("Ratings Distribution")
plt.xlabel("Rating")
plt.ylabel("Count")
plt.show()
```




```
In [14]: # rating timestamp analysis;if ratings show clear temporal trends, keep transfor
if 'rating_timestamp' in df_merged.columns:
    df_merged['rating_timestamp'] = pd.to_datetime(df_merged['rating_timestamp'])
    df_merged['rating_year'] = df_merged['rating_timestamp'].dt.year

    print("\n=== Ratings per Year ===")
    print(df_merged['rating_year'].value_counts().sort_index().head())
```

```
df_merged['rating_year'].value_counts().sort_index().plot(kind="bar", figsize=
plt.title("Ratings Over Time")
plt.xlabel("Year")
plt.ylabel("Count")
plt.show()
```

```
=== Ratings per Year ===
rating_year
1996      6040
1997      1916
1998       507
1999      2439
2000     10074
Name: count, dtype: int64
```



```
In [15]: # Tag analysis; if too sparse, may not be useful
if 'tag' in df_merged.columns:
    print("\n=== Most Common Tags ===")
    print(df_merged['tag'].value_counts().head(10))

    tag_missing_ratio = df_merged['tag'].isnull().mean()
    print(f"Tag missing ratio: {tag_missing_ratio:.2f}")
```

```
=== Most Common Tags ===
tag
In Netflix queue      55
atmospheric           36
Disney                23
superhero             23
funny                 23
surreal               23
religion              22
thought-provoking     22
quirky                21
psychology            20
Name: count, dtype: int64
Tag missing ratio: 0.97
```

```
In [16]: # Genre analysis
if 'genres' in df_merged.columns:
    print("\n=== Most Common Genres ===")
    print(df_merged['genres'].value_counts().head(10))
```

```
=== Most Common Genres ===
genres
Comedy                7250
Drama                 6403
Comedy|Romance        4001
Comedy|Drama|Romance  3039
Drama|Romance         2880
Comedy|Drama          2863
Action|Adventure|Sci-Fi 2430
Crime|Drama           2347
Action|Crime|Thriller  1574
Action|Adventure|Sci-Fi|Thriller 1463
Name: count, dtype: int64
```

```
In [17]: # External IDs analysis
if 'imdbId' in df_merged.columns and 'tmdbId' in df_merged.columns:
    print("\n=== External IDs Info ===")
    print("Unique imdbIds:", df_merged['imdbId'].nunique())
    print("Unique tmdbIds:", df_merged['tmdbId'].nunique())
```

```
=== External IDs Info ===
Unique imdbIds: 9724
Unique tmdbIds: 9715
```

```
In [18]: # Final dataframe selection
Df_columns_keep = ['userId', 'movieId', 'rating', 'rating_year', 'genres', 'title']
df_final = df_merged[Df_columns_keep]
df_final.head()
```

```
Out[18]:
```

	userId	movieId	rating	rating_year	genres	
0	1	1	4.0	2000	Adventure Animation Children Comedy Fantasy	Toy Story
1	1	3	4.0	2000	Comedy Romance	Grumpier C
2	1	6	4.0	2000	Action Crime Thriller	Heat
3	1	47	5.0	2000	Mystery Thriller	Seven (a.k.a
4	1	50	5.0	2000	Crime Mystery Thriller	Usual Suspe

Data Preprocessing

```
In [19]: # Encode categorical features
user_enc = LabelEncoder()
movie_enc = LabelEncoder()
genre_enc = LabelEncoder()
title_enc = LabelEncoder()

df_final['userId_enc'] = user_enc.fit_transform(df_final['userId'])
df_final['movieId_enc'] = movie_enc.fit_transform(df_final['movieId'])
df_final['genres_enc'] = genre_enc.fit_transform(df_final['genres'].astype(str))
df_final['title_enc'] = title_enc.fit_transform(df_final['title'])
```

```
In [20]: # Features (X) and Target (y)
X = df_final[['userId_enc', 'movieId_enc', 'rating_year', 'genres_enc', 'title_en
y = df_final['rating']
```

```
In [21]: # Split ratings data into train and test sets
train_ratings, test_ratings = train_test_split(df_final, test_size=0.2, random_s

# Create user-movie matrix
user_movie_matrix_train = train_ratings.pivot_table(index='userId', columns='mov
user_movie_matrix_test = test_ratings.pivot_table(index='userId', columns='movie

# Align dimensions; if test has unseen users/movies
user_movie_matrix_test = user_movie_matrix_test.reindex(index=range(n_users), co
user_movie_matrix_train = user_movie_matrix_train.reindex(index=range(n_users), c

X_train = user_movie_matrix_train.values
y_train = user_movie_matrix_train.values
```

```
In [33]: # Display shape and a sample
print("Matrix shape:", user_movie_matrix_train.shape)
user_movie_matrix_train.head()
```

Matrix shape: (610, 9724)

```
Out[33]: movieId    0    1    2    3    4    5    6    7    8    9    ...  9714  9715  9716  9717  9718  9
          userId
0      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...   0.0   0.0   0.0   0.0   0.0
1      0.0  4.0  0.0  4.0  0.0  0.0  4.0  0.0  0.0  0.0  ...   0.0   0.0   0.0   0.0   0.0
2      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...   0.0   0.0   0.0   0.0   0.0
3      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...   0.0   0.0   0.0   0.0   0.0
4      0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...   0.0   0.0   0.0   0.0   0.0
```

5 rows × 9724 columns

```
In [35]: # Compute user-user similarity matrix
user_similarity = cosine_similarity(user_movie_matrix_train)
user_similarity_df = pd.DataFrame(
    user_similarity,
    index=user_movie_matrix_train.index,
    columns=user_movie_matrix_train.index
)
# Display a sample
user_similarity_df.head()
```

```
Out[35]:
```

userId	0	1	2	3	4	5	6	7	8
0	0.0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
1	0.0	1.000000	0.034755	0.039183	0.165461	0.137692	0.124770	0.143811	0.135497
2	0.0	0.034755	1.000000	0.000000	0.000000	0.041284	0.064996	0.068174	0.000000
3	0.0	0.039183	0.000000	1.000000	0.002961	0.006385	0.003619	0.000000	0.006890
4	0.0	0.165461	0.000000	0.002961	1.000000	0.130157	0.085226	0.125647	0.048075

5 rows × 610 columns

Define pipeline and Gridsearchcv

```
In [23]: # Define pipeline
pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="constant", fill_value=0)),
    ('scaler', StandardScaler()),
    ("svd", TruncatedSVD()),
    ("ridge", Ridge()),
])

# Hyperparameter grid
param_grid = {
    "svd__n_components": [20, 50, 100],
    "ridge__alpha": [0.01, 0.1, 1, 10]
}

# Cross-validation
cv = KFold(n_splits=5, shuffle=True, random_state=42)

grid = GridSearchCV(pipe, param_grid, cv=cv, scoring="neg_mean_squared_error", v
    # Fit on training matrix (users x movies)
    grid.fit(X_train, y_train)

print("Best Params:", grid.best_params_)
print("Best CV Score (MSE):", -grid.best_score_)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

Best Params: {'ridge__alpha': 1, 'svd__n_components': 100}

Best CV Score (MSE): 0.10228348384991699

Make prediction and Model evaluation

```
In [25]: # Predictions for test users
preds = grid.best_estimator_.predict(user_movie_matrix_test.values)

# Compute RMSE
y_true, y_pred = [], []
for row in range(user_movie_matrix_test.shape[0]):
    for col in range(user_movie_matrix_test.shape[1]):
        if user_movie_matrix_test.iloc[row, col] > 0: # only evaluate on actual
            y_true.append(user_movie_matrix_test.iloc[row, col])
            y_pred.append(preds[row, col])

rmse = np.sqrt(mean_squared_error(y_true, y_pred))
print("Test RMSE:", rmse)
```


Test RMSE: 3.2944222855896497

Recommend top 5 movies for a user

```
In [31]: def recommend_movies(user_id, preds, df, top_n=5):
    user_idx = user_enc.transform([user_id])[0]
    user_preds = preds[user_idx]

    # Movies already rated by the user
    rated_movies = df[df['userId'] == user_id]['movieId'].values
    rated_items = movie_enc.transform(rated_movies)

    # Exclude rated movies
    user_preds = user_preds.copy()
    user_preds[rated_items] = -np.inf

    # Top-N recommendations
    top_items = np.argsort(user_preds)[-top_n:][::-1]
    movie_ids = movie_enc.inverse_transform(top_items)
    return df[df['movieId'].isin(movie_ids)][['title', 'genres']].drop_duplicates()

# Example: recommend for user 1
print(recommend_movies(1, preds, df_final, top_n=5))
```

	title	genres
735	I Love Trouble (1994)	Action Comedy
753	Age of Innocence, The (1993)	Drama
2347	Virtuosity (1995)	Action Sci-Fi Thriller
14411	Cemetery Man (Dellamorte Dellamore) (1994)	Horror
86979	Band of the Hand (1986)	Action Crime Drama

Conclusion and Recommendation

Our recommender system shows that collaborative filtering with matrix factorization can generate and personalized movie suggestions. By using user ratings and key metadata, the model achieves accuracy with an RMSE of 3.29 meaning most recommendations matched user interests. However, to improve the model, future work should explore hybrid approaches that combine collaborative and content-based filtering, add time-awareness to capture changing preferences, and focus on increasing diversity in recommendations. Overall, this project provides a strong baseline for a movie recommendation engine and a foundation for building more advanced systems.